


MOCKUP — Tenant Fairness with vLLM Continuous Batching

Inference requests execute concurrently; fairness is enforced when batch slots become available

Both tenants are model-ready and continuously backlogged; four batch slots; each request generates four decode tokens. Dashed lines mark new admissions.

Tenant A request Tenant B request

Preparation-only: native/FCFS inference admission
A = 36 slot-steps, B = 12 slot-steps; service gap = 24



Cross-stage max-min: tenant-fair inference admission
A = 24 slot-steps, B = 24 slot-steps; service gap = 0

