


Platform Service Profiler: Preparation and Inference Cost

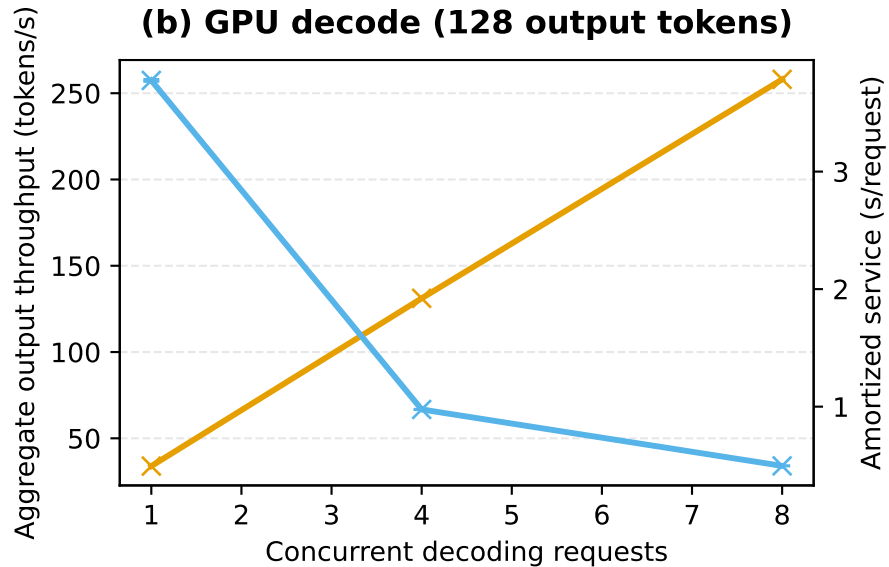
✕ Throughput/capacity

✕ Profiled service cost

(a) CPU preparation



(b) GPU decode (128 output tokens)



Measured on the A40 platform; points are profile means/medians and error bars show variation across runs.