


MOCKUP 1 — Constant-Rate Two-Tenant Service Fairness

Heterogeneous CPU and inference costs are measured in separate matched bottleneck runs

FCFS Tenant RR Prep-only Inference-only Cross-stage

CPU run: equal rates, A=128 frames and B=32. Inference run: equal rates and 32 frames, A requests ~100 generated tokens and B one word. Both tenants remain stage-backlogged.

(a) Preparation-heavy configuration



(b) Inference-heavy configuration

