


MOCKUP 2 — ON/OFF Work Conservation Under Cross-Stage Max-Min

Preparation and inference are evaluated in separate capacity-calibrated runs

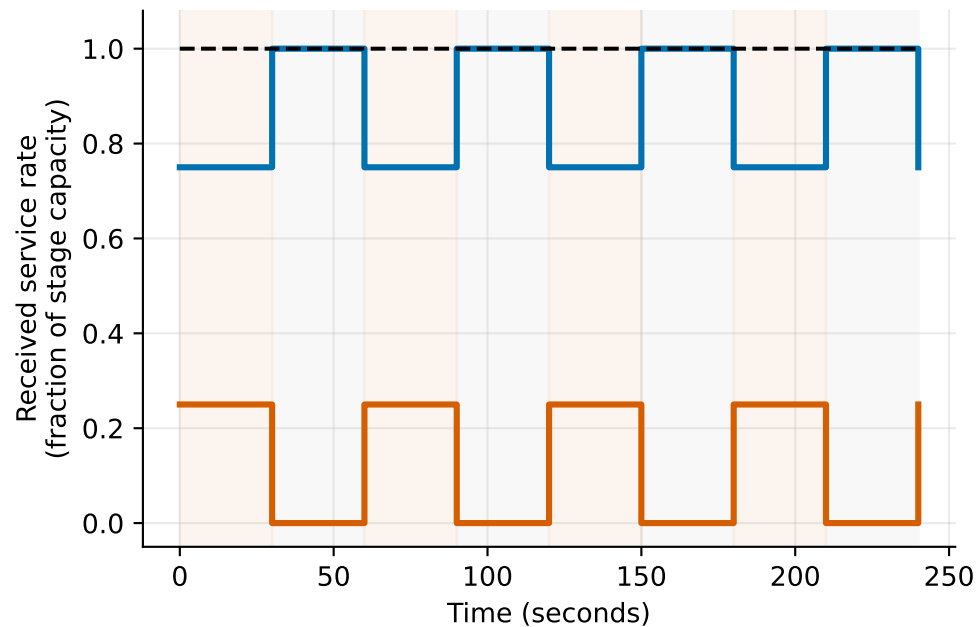
Two matched runs with empirically calibrated stage capacity C ; identical 32-frame requests.
A offers $0.25C$ for 30 s ON, then zero for 30 s OFF; B remains backlogged. Use 30 s windows.

A ON Tenant A received --- Total utilized capacity
A OFF Tenant B received

(a) Preparation service rate



(b) Inference service rate



(c) A is protected; B's overload backlog grows

