


CPU media preparation GPU visual prefill GPU output decoding



Lower service cost per request
⇒ higher request throughput

Equal request counts do not imply equal service across CPU and GPU stages.